

Importing Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn import svm
from sklearn.tree import DecisionTreeRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
```

Loading Data

```
In [3]: data = pd.read_csv(r'C:\Users\vamsi\Desktop\M.Tech\ML\19 Projects\medical_cost_da
```

```
In [4]: data.head()
```

Out[4]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [5]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [6]: data.describe()
```

```
Out[6]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [7]: data.isnull().sum()
```

```
Out[7]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

Categorical data

```
In [8]: data.sex.unique()
```

```
Out[8]: array(['female', 'male'], dtype=object)
```

```
In [9]: data.smoker.unique()
```

```
Out[9]: array(['yes', 'no'], dtype=object)
```

```
In [10]: data.region.unique()
```

```
Out[10]: array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)
```

Label Encoding

```
In [11]: from sklearn.preprocessing import LabelEncoder
#sex
le = LabelEncoder()
le.fit(data.sex.drop_duplicates())
data.sex = le.transform(data.sex)

# smoker or not
le.fit(data.smoker.drop_duplicates())
data.smoker = le.transform(data.smoker)

#region
le.fit(data.region.drop_duplicates())
data.region = le.transform(data.region)
```


Data Correlation

```
In [12]: data.corr()['charges'].sort_values()
```

```
Out[12]: region      -0.006208
sex              0.057292
children         0.067998
bmi              0.198341
age              0.299008
smoker           0.787251
charges          1.000000
Name: charges, dtype: float64
```

```
In [13]: f, ax = plt.subplots(figsize=(10, 8))  
corr = data.corr()  
sns.heatmap(corr, ax =ax, annot = True)
```

Out[13]: <AxesSubplot:>




Clearly, a strong correlation is observed only with the fact of smoking the patient

Smoker Analysis

```
In [14]: sns.factorplot(x="smoker", kind="count", hue = 'sex', data=data)
```

C:\Users\vamsi\anaconda3\lib\site-packages\seaborn\categorical.py:3704: UserWarning: The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed to `strip` in `catplot`.
warnings.warn(msg)

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x1ee05150520>
```



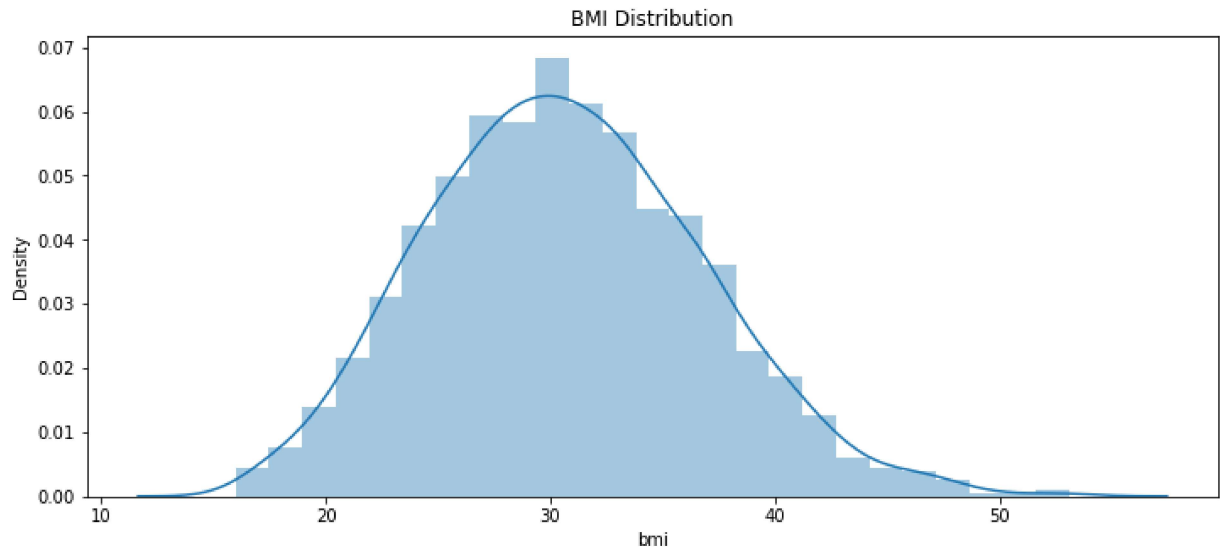
The total cost of treatment in men will be more than in women, given the impact of smoking.

BMI Analysis

```
In [15]: plt.figure(figsize=(12,5))
plt.title("BMI Distribution")
ax = sns.distplot(data["bmi"])
```

C:\Users\vamsi\anaconda3\lib\site-packages\seaborn\distributions.py:2551: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)



The average BMI in patients is 30.

Children Count Analysis

```
In [16]: sns.factorplot(x="children", kind="count", data=data, size = 6)
```


C:\Users\vamsi\anaconda3\lib\site-packages\seaborn\categorical.py:3704: UserWarning: The `factorplot` function has been renamed to `catplot`. The original name will be removed in a future release. Please update your code. Note that the default `kind` in `factorplot` (`'point'`) has changed to `strip` in `catplot`.

warnings.warn(msg)

C:\Users\vamsi\anaconda3\lib\site-packages\seaborn\categorical.py:3710: UserWarning: The `size` parameter has been renamed to `height`; please update your code.

warnings.warn(msg, UserWarning)

```
Out[16]: <seaborn.axisgrid.FacetGrid at 0x1ee0adb8370>
```



Most patients do not have children.

In [17]: data.head()

Out[17]:

	age	sex	bmi	children	smoker	region	charges
0	19	0	27.900	0	1	3	16884.92400
1	18	1	33.770	1	0	2	1725.55230
2	28	1	33.000	3	0	2	4449.46200
3	33	1	22.705	0	0	1	21984.47061
4	32	1	28.880	0	0	1	3866.85520

Splitting Data

In [18]: *# training dataset and labels*
x = data.drop(data.columns[[6, 5]], axis = 1)
y = data['charges']

In [19]: x.head()

Out[19]:

	age	sex	bmi	children	smoker
0	19	0	27.900	0	1
1	18	1	33.770	1	0
2	28	1	33.000	3	0
3	33	1	22.705	0	0
4	32	1	28.880	0	0

In [21]: *# splitting into training and testing data*
from sklearn.model_selection import train_test_split
xtrain, xtest, ytrain, ytest = train_test_split(x, y, test_size = 0.30, random_st

Feature Scaling

In [22]: from sklearn.preprocessing import MinMaxScaler
sc_x = MinMaxScaler()
xtrain = sc_x.fit_transform(xtrain)
xtest = sc_x.fit_transform(xtest)

In [23]: xtrain[0:2]

Out[23]: array([[0. , 0. , 0.32970137, 0. , 0.],
[0.45652174, 0. , 0.45305354, 0. , 0.]])

ML Models


```
In [24]: linear = LinearRegression()  
dt = DecisionTreeRegressor()  
svr = svm.SVR()  
rf = RandomForestRegressor()
```

Training

```
In [25]: linear.fit(xtrain, ytrain)
```

```
Out[25]: LinearRegression()
```

```
In [26]: dt.fit(xtrain, ytrain)
```

```
Out[26]: DecisionTreeRegressor()
```

```
In [27]: svr.fit(xtrain, ytrain)
```

```
Out[27]: SVR()
```

```
In [28]: rf.fit(xtrain, ytrain)
```

```
Out[28]: RandomForestRegressor()
```

Prediction

```
In [29]: y_pred_linear = linear.predict(xtest)  
y_pred_dt = dt.predict(xtest)  
y_pred_svr = svr.predict(xtest)  
y_pred_rf = rf.predict(xtest)
```

RMSE Error

```
In [30]: import math  
error_linear = math.sqrt(mean_squared_error(y_pred_linear, ytest))  
error_dt = math.sqrt(mean_squared_error(y_pred_dt, ytest))  
error_svr = math.sqrt(mean_squared_error(y_pred_svr, ytest))  
error_rf = math.sqrt(mean_squared_error(y_pred_rf, ytest))
```

```
In [31]: print ("      Model          :      RMSE Error\n")
print ("Linear Regression   : ", error_linear)
print ("Decision Tree       : ", error_dt)
print ("Support Vector       : ", error_svr)
print ("Random Forest        : ", error_rf)
```

Model	:	RMSE Error
Linear Regression	:	5827.6103488260305
Decision Tree	:	7291.157424644265
Support Vector	:	13132.937356154958
Random Forest	:	5164.067609731632

```
In [ ]:
```